

## xgboost분류.py

```
1 import pandas as pd
2 import numpy as np
3 from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
4 from sklearn.preprocessing import LabelEncoder
5 from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score
6 import xgboost as xgb
7
8 # Step 1: Load the Titanic dataset
9 # Assuming the data is downloaded from Kaggle and placed in the working directory
10 train_df = pd.read_csv('train.csv')
11 test_df = pd.read_csv('test.csv')
12
13 # Step 2: Data Preprocessing
14 # Handle missing values
15 train_df['Age'].fillna(train_df['Age'].median(), inplace=True)
16 test_df['Age'].fillna(test_df['Age'].median(), inplace=True)
17 train_df['Embarked'].fillna(train_df['Embarked'].mode()[0], inplace=True)
18 test_df['Fare'].fillna(test_df['Fare'].median(), inplace=True)
19
20 # Drop unnecessary columns
21 train_df.drop(['Cabin', 'Ticket', 'Name', 'PassengerId'], axis=1, inplace=True)
22 test_passenger_ids = test_df['PassengerId'] # Save for submission
23 test_df.drop(['Cabin', 'Ticket', 'Name', 'PassengerId'], axis=1, inplace=True)
24
25 # Encode categorical variables
26 label_encoder = LabelEncoder()
27 train_df['Sex'] = label_encoder.fit_transform(train_df['Sex'])
28 test_df['Sex'] = label_encoder.transform(test_df['Sex'])
29
30 train_df['Embarked'] = label_encoder.fit_transform(train_df['Embarked'])
31 test_df['Embarked'] = label_encoder.transform(test_df['Embarked'])
32
33 # Split train data into features and target
34 X = train_df.drop('Survived', axis=1)
35 y = train_df['Survived']
36
37 # Split into train and validation sets
38 X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=42)
39
40 # Step 3: Define XGBoost model
41 model = xgb.XGBClassifier(objective='binary:logistic', random_state=42)
42
43 # Step 4: Parameter optimization using GridSearchCV (which includes cross-validation)
44 param_grid = {
45     'n_estimators': [50, 100, 200],
46     'max_depth': [3, 5, 7],
47     'learning_rate': [0.01, 0.1, 0.2],
48     'subsample': [0.8, 1.0],
49     'colsample_bytree': [0.8, 1.0]
50 }
51
```

```

52 grid_search = GridSearchCV(estimator=model, param_grid=param_grid, cv=5, scoring='accuracy',
53                             n_jobs=-1, verbose=1)
54
55 # Best parameters and model
56 best_params = grid_search.best_params_
57 print("Best parameters found: ", best_params)
58 best_model = grid_search.best_estimator_
59
60 # Step 5: Evaluate model performance on validation set
61 y_pred_val = best_model.predict(X_val)
62
63 # Performance metrics
64 accuracy = accuracy_score(y_val, y_pred_val)
65 precision = precision_score(y_val, y_pred_val)
66 recall = recall_score(y_val, y_pred_val)
67 f1 = f1_score(y_val, y_pred_val)
68 roc_auc = roc_auc_score(y_val, y_pred_val)
69
70 print("\nModel Performance on Validation Set:")
71 print(f"Accuracy: {accuracy:.4f}")
72 print(f"Precision: {precision:.4f}")
73 print(f"Recall: {recall:.4f}")
74 print(f"F1 Score: {f1:.4f}")
75 print(f"ROC AUC: {roc_auc:.4f}")
76
77 # Cross-validation scores on full train data
78 cv_scores = cross_val_score(best_model, X, y, cv=5, scoring='accuracy')
79 print("\nCross-Validation Accuracy Scores: ", cv_scores)
80 print("Mean CV Accuracy: ", cv_scores.mean())
81
82 # Step 6: Train the best model on full train data
83 best_model.fit(X, y)
84
85 # Step 7: Predict on test data
86 test_predictions = best_model.predict(test_df)
87
88 # Create submission file
89 submission = pd.DataFrame({
90     'PassengerId': test_passenger_ids,
91     'Survived': test_predictions
92 })
93 submission.to_csv('submission.csv', index=False)
94 print("\nPredictions saved to submission.csv")

```